

What Does This Strategy Actually Return Over a Real, Continuous 6-Month Stretch?

6 independent, continuous, regime-spanning 6-month windows · current production settings, completely unchanged · no cherry-picking

Generated 2026-06-21 · BTC/ETH/SOL · 15m + 1h · self-learning active throughout

Every prior result — the original Model Validation Report, the held-out ADA/XRP test, and the Addendum's two follow-up batches — measured short windows (30-90 days), scattered random spot-checks, or isolated pattern/timeframe questions. None of them directly answered "if I ran this, unmodified, for 6 continuous months, what would I actually see?" This campaign produces that number: 6 independent \$100,000 runs, each one unbroken 6-month historical stretch, deliberately spanning known bull, bear, and choppy regimes — with the current production configuration (1% risk, 1.5 R:R, 65 confidence floor, all guardrails, self-learning hints active and accumulating) left completely untouched.

Configuration & Window Selection

SETTING	VALUE
Symbols	BTCUSDT, ETHUSDT, SOLUSDT
Timeframes	15m, 1h (5m and 4h excluded — see Addendum)
Window length	6 continuous months, no gaps, no stitching
Starting capital	\$100,000 per window (independent — windows do not chain)
Risk per trade / Min R:R / Min confidence / Score floor	1.0% / 1.5 / 65 / 50 (unchanged)
Regime filter / Session filter / HTF confirmation	ON / ON / ON (unchanged)
Guardrails (daily 3%, weekly 6%, cooldown after 3 losses, drawdown halt 50%)	All unchanged
Self-learning (db.edgeStats)	Active and accumulating throughout — not frozen

One deviation from the as-specified defaults, applied deliberately: `sessionFilter` defaults to **OFF** in `lib/simulator.ts`'s `defaultSimConfig()` — it was off in every prior batch in this entire project (original report, held-out test, Batch A, Batch B). The campaign brief explicitly calls for "Session filter: ON," so it was turned on for this campaign only, via the existing config flag. This is flagged here rather than silently applied, since it makes this campaign not directly apples-to-apples with every prior batch on that one dimension — though session filtering only restricts *when* a trade can open, not the underlying signal logic.

Window selection — one substitution, explicitly noted

#	HAND-PICKED REGIME	DATE RANGE	REASONING
1	Bull	2020-10-01 → 2021-03-31	Run-up to the first 2021 ATH
2	Bull	2023-10-01 → 2024-03-31	Run-up into the March 2024 ATH
3	Bear	2021-05-01 → 2021-10-31	China mining-ban crash + chop recovery
4	Bear	2022-04-01 → 2022-09-30	Terra/Luna collapse + broader 2022 bear market
5	Choppy	2023-07-01 → 2023-12-31	Pre-ETF-approval grind, documented low-volatility chop
6	Choppy (substitute)	2024-04-01 → 2024-09-30	Originally specified as Sep 2019 - Feb 2020, but SOLUSDT has no Binance trading history before 2020-08-11 — that window is impossible for one of the three required symbols. Substituted the well-documented post-March-2024-ATH consolidation/chop period, which immediately follows Window 2 with no date overlap.

Worth noting, not corrected: Window 2 (Oct 2023 - Mar 2024, bull) and Window 5 (Jul 2023 - Dec 2023, choppy) overlap by three months (Oct-Dec 2023) — both were specified explicitly by name in the campaign brief with independent reasoning, so the overlap was kept rather than second-guessed, but it means these two windows are not fully independent measurements; roughly half of Window 5's months are also inside Window 2.

Per-Window Results

#	DATE RANGE	HAND-PICKED	CLASSIFIER'S LABEL	TRADES	WIN RATE	AVG PF	AVG RETURN	AVG MAX DD
1	2020-10 → 2021-03	Bull	weak_uptrend (agrees)	653	55.8%	1.31	+20.5%	12.2%
2	2023-10 → 2024-03	Bull	weak_uptrend (agrees)	813	54.3%	1.30	+22.6%	17.6%
3	2021-05 → 2021-10	Bear	weak_uptrend (disagrees)	719	56.1%	1.22	+16.0%	17.8%
4	2022-04 → 2022-09	Bear	weak_downtrend (agrees)	818	54.1%	1.28	+27.2%	16.0%
5	2023-07 → 2023-12	Choppy	ranging (agrees)	770	46.5%	0.97	-6.1%	19.1%
6	2024-04 → 2024-09	Choppy (sub.)	weak_downtrend (partial)	789	56.2%	1.64	+43.0%	13.1%

Classifier disagreement, flagged as instructed: Window 3 (May-Oct 2021) is unambiguously a known bear/crash period — the China mining-ban crash took BTC from ~\$58k to ~\$29k by July 2021 — yet `classifyRegime()` 's per-run dominant label across this window's 6 sub-runs was **weak_uptrend**, not weak/strong downtrend. The likely explanation: the same 6-month window also contains the sharp Jul-Oct 2021 recovery rally back toward new highs, and `classifyRegime()` only samples while the strategy is flat (not mid-trade) using a trailing 50-candle EMA20/50 window — if the recovery phase produced more "flat" candles than the initial crash phase, it would dominate the per-run tally even though the window *as a whole* is correctly remembered as a bear/crash period by anyone who lived through it. This isn't necessarily a classifier bug; it's a reminder that "dominant regime over 6 months" is a genuinely ambiguous question when a single window contains a sharp regime change partway through — worth a closer look if regime-conditional tuning is ever pursued, but not addressed here per scope.

Distribution Summary

+20.5% MEAN RETURN	+21.5% MEDIAN RETURN	14.6% STD. DEVIATION	16.0% MEAN MAX DRAWDOWN
--	--	--	---

- **Best window:** #6, Choppy (substitute) Apr-Sep 2024, **+43.0%**
- **Worst window:** #5, Choppy Jul-Dec 2023, **-6.1%**
- **Worst single sub-run drawdown:** 32.1% (ETHUSDT 15m, Window 6) — well short of the 50% drawdown-halt threshold; the halt never triggered anywhere in this campaign
- **Windows reaching ≥30% return:** 1 of 6 (Window 6)
- **Windows net negative:** 1 of 6 (Window 5)

A standard deviation of 14.6% against a mean of 20.5% is substantial — roughly 70% of the spread of the mean itself. The two choppy windows alone produced the single best (+43.0%) and single worst (-6.1%) results in the entire campaign. **This is a genuinely bimodal result for the choppy regime, not a moderate one** — averaging it into one number ("choppy: +18.4% mean") would misrepresent what's actually happening, which is "choppy can go either way, hard," not "choppy is fine."

By Hand-Picked Regime

REGIME	WINDOWS	MEAN RETURN	MEAN MAX DD
Bull	#1 (+20.5%), #2 (+22.6%)	+21.5%	14.9%
Bear	#3 (+16.0%), #4 (+27.2%)	+21.6%	16.9%
Choppy	#5 (-6.1%), #6 (+43.0%)	+18.4%	16.1%

Bull and bear produced almost identical mean returns (21.5% vs 21.6%) — this strategy does not show a strong directional bias toward trending-up conditions despite HTF confirmation existing specifically to ride trends. It seems equally capable of catching bear-market relief rallies (Window 4's Terra/Luna aftermath: +27.2%) as bull runs. Choppy is the outlier — not because its mean is bad (+18.4%, actually fine), but because its *variance* is by far the highest of the three regimes, swinging from the campaign's single best result to its single worst.

By Timeframe — the most actionable single finding in this campaign

TIMEFRAME	MEAN RETURN	MEAN PF	MEAN WR	PROFITABLE SUB-RUNS	WINDOWS ≥ 30%
15m	+14.2%	1.07	48.9%	9/18 (50%)	—
1h	+26.9%	1.50	58.7%	16/18 (89%)	3/6

Restricted to 1h alone, the mean return across the same 6 windows is **26.9%** — close to the 30% target — with **89% of sub-runs profitable** and 3 of 6 windows individually clearing 30% (Window 2: 30.2%, Window 4: 35.5%, Window 6: 56.8%). 15m drags the combined average down: 50% of its sub-runs were profitable, and several were sharply negative (BTCUSDT 15m hit -23.9% in Window 2 and -23.8% in Window 3). This mirrors every prior finding in this project's history — 1h has consistently outperformed 15m, which has consistently outperformed 5m — but this is now backed by continuous 6-month evidence rather than 30-90 day windows.

Guardrail Activity

WINDOW	DAILY/WEEKLY LIMIT REJECTIONS*	LOSS-STREAK COOLDOWN (CANDLES, ~EPISODES)	MAX DRAWDOWN SEEN	DRAWDOWN HALT
1	1,874	840 (~42)	17.9%	No
2	1,108	860 (~43)	25.7%	No
3	986	751 (~38)	27.4%	No
4	1,876	1,152 (~58)	26.0%	No
5	1,420	1,040 (~52)	26.0%	No
6	1,572	874 (~44)	32.1%	No

* `tradesRejectedByDailyLimit` is a shared counter for both the daily (3%) and weekly (6%) loss limits in the underlying code — they cannot be split apart from the data captured here. Episode counts for the cooldown are estimated as candles-rejected ÷ 20 (the fixed cooldown length), since the engine tracks candles skipped, not discrete trigger events; overlapping triggers would understate the true episode count slightly.

No drawdown halt fired anywhere in this campaign — the worst single sub-run drawdown (32.1%) stayed comfortably under the 50% halt threshold. **Guardrail activity was broadly similar across all 6 windows** — no

single regime triggered dramatically more daily/weekly limit hits or cooldowns than the others (Window 4, the Terra/Luna bear market, shows the highest cooldown-episode estimate at ~58, which is plausible given that was the most acutely stressful period tested, but it's not a dramatic outlier against the other windows' 38-52 range). There is no "guardrails protected window X from being far worse" story to tell here distinctly from any other window — they were active throughout, fairly evenly.

Explicit Verdict

Across 6 continuous 6-month windows spanning bull, bear, and choppy regimes, the current production configuration averaged +20.5% return (median +21.5%) with a mean max drawdown of 16.0%. 1 of 6 windows reached the 30% target. The strategy is not yet consistently approaching the 30%-in-6-months goal under current settings, combining 15m and 1h together. It performs almost identically in bull and bear conditions (+21.5% vs +21.6% mean) and is highly inconsistent in choppy conditions specifically (-6.1% to +43.0%, the widest spread of any regime). The single clearest lever available: **1h alone averages +26.9% with 89% of sub-runs profitable**, far closer to target than 15m's +14.2% with only 50% profitable — most of the shortfall against the 30% goal is attributable to 15m underperformance, not a uniform weakness across both timeframes.

Per the decision framework this campaign was run to inform: mean return is meaningfully below 30%, so the next step is Phase 3 — testing the four identified levers (score floor, conviction-weighted sizing, exit logic, and conditionally 5m-Hammer-only if it separately clears PF 1.0), each re-run against these same 6 fixed windows so the before/after comparison stays clean. Given what this campaign found, the highest-leverage starting point within Phase 3 is likely whatever can close the 15m gap specifically — score floor and exit-logic tuning seem more directly relevant to that than sizing, since 15m's win rate (48.9%) and PF (1.07) are both already only marginally above breakeven, suggesting an entry-quality problem on that timeframe rather than a sizing problem.

Appendix — All 36 Sub-Run Results

WINDOW	SYMBOL	TF	TRADES	WR	PF	RETURN	MAX DD
1	BTCUSDT	15m	172	48.8%	1.08	+10.3%	17.9%
1	BTCUSDT	1h	68	67.6%	2.03	+54.2%	9.4%
1	ETHUSDT	15m	179	54.2%	0.97	-3.3%	13.1%
1	ETHUSDT	1h	56	57.1%	1.26	+11.3%	16.0%
1	SOLUSDT	15m	146	56.8%	1.36	+45.4%	10.2%
1	SOLUSDT	1h	32	50.0%	1.15	+4.9%	6.7%
2	BTCUSDT	15m	220	45.0%	0.79	-23.9%	24.6%
2	BTCUSDT	1h	73	63.0%	1.93	+53.3%	12.9%
2	ETHUSDT	15m	181	47.5%	0.95	-6.2%	25.7%
2	ETHUSDT	1h	91	53.8%	1.38	+26.4%	11.0%
2	SOLUSDT	15m	199	53.3%	1.46	+74.9%	16.8%
2	SOLUSDT	1h	49	63.3%	1.31	+11.0%	14.3%
3	BTCUSDT	15m	204	45.1%	0.78	-23.8%	25.8%
3	BTCUSDT	1h	86	52.3%	0.92	-6.6%	21.1%
3	ETHUSDT	15m	169	53.3%	1.31	+49.7%	27.4%
3	ETHUSDT	1h	63	63.5%	1.44	+20.5%	13.6%
3	SOLUSDT	15m	154	61.7%	1.32	+36.7%	12.4%
3	SOLUSDT	1h	43	60.5%	1.56	+19.4%	6.4%

4	BTCUSDT	15m	198	40.9%	0.86	-16.7%	26.0%
4	BTCUSDT	1h	68	57.4%	1.23	+12.0%	9.9%
4	ETHUSDT	15m	193	51.8%	1.32	+43.5%	20.4%
4	ETHUSDT	1h	86	64.7%	1.87	+78.5%	11.7%
4	SOLUSDT	15m	206	51.5%	1.16	+30.1%	16.7%
4	SOLUSDT	1h	67	58.2%	1.26	+16.0%	11.3%
5	BTCUSDT	15m	193	42.5%	0.94	-8.8%	23.2%
5	BTCUSDT	1h	81	47.5%	0.81	-12.2%	20.9%
5	ETHUSDT	15m	153	41.2%	0.78	-21.3%	26.0%
5	ETHUSDT	1h	70	46.4%	1.03	+1.7%	10.5%
5	SOLUSDT	15m	200	42.5%	0.86	-18.8%	23.5%
5	SOLUSDT	1h	73	58.9%	1.38	+22.6%	10.7%
6	BTCUSDT	15m	192	47.9%	1.05	+6.2%	15.2%
6	BTCUSDT	1h	65	64.6%	2.55	+80.9%	6.9%
6	ETHUSDT	15m	190	40.0%	0.61	-31.6%	32.1%
6	ETHUSDT	1h	76	56.6%	1.26	+14.9%	7.1%
6	SOLUSDT	15m	203	56.9%	1.71	+112.6%	10.0%
6	SOLUSDT	1h	63	71.0%	2.62	+74.7%	6.9%

Report generated from 36 individual live simulation runs against real historical Binance market data, captured directly from each run's full in-memory result object (including regime breakdown and rejection counters not persisted to the database). Zero console errors across the entire campaign. Raw data: `training-exports/baseline-campaign-results.json`.